

Lab session 5

Discrete cosine transform Introduction to JPEG

report by Grebennikova Anna & Kseniia Ovchinnikova

Exercise 1. Discrete cosine transform

Let f be a discrete signal defined on $\{0, \dots, L-1\}$, consider the signal \tilde{f} built by symmetrizing f with respect to $-\frac{1}{2}$, namely

$$\tilde{f}_n = \begin{cases} f_n & \text{for } 0 \leq n < L, \\ f_{-n-1} & \text{for } -L \leq n \leq -1. \end{cases}$$

We then assume \tilde{f} is periodized with period $2L$.

Question 1.1

We aim to prove that f_n admits the following decomposition:

$$f_n = \sum_{k=0}^{L-1} a_k \cos\left(\frac{k\pi}{L} \left(n + \frac{1}{2}\right)\right)$$

To establish this, let us calculate f_{-n-1} :

$$f_{-n-1} = \sum_{k=0}^{L-1} a_k \cos\left(\frac{k\pi}{L} \left(-n-1 + \frac{1}{2}\right)\right) = \sum_{k=0}^{L-1} a_k \cos\left(\frac{k\pi}{L} \left(-n - \frac{1}{2}\right)\right) = \sum_{k=0}^{L-1} a_k \cos\left(\frac{k\pi}{L} \left(n + \frac{1}{2}\right)\right) = f_n$$

Therefore, we conclude:

$$\Rightarrow f_n = f_{-n-1}$$

This result demonstrates that f_n possesses a symmetry property, namely $f_n = f_{-n-1}$. Consequently, the function f_n is symmetric about $-\frac{1}{2}$, and the given decomposition holds true. The decomposition effectively captures the periodic and symmetric nature of the discrete signal f_n .

To determine the coefficients a_k , we use the following approach:

$$f_n = \frac{\tilde{f}_n + \tilde{f}_{-n-1}}{2}$$

We know that

$$\tilde{f}_n = \frac{1}{2L} \sum_{k=0}^{2L-1} \hat{f}_k e^{\frac{2i\pi kn}{2L}}$$

Substituting this expression for \tilde{f}_n and \tilde{f}_{-n-1} into the formula for f_n , we obtain:

$$\begin{aligned} f_n &= \frac{1}{4L} \left(\sum_{k=0}^{2L-1} \hat{f}_k e^{\frac{2i\pi kn}{2L}} + \sum_{k=0}^{2L-1} \hat{f}_k e^{\frac{2i\pi k(-n-1)}{2L}} \right) = \frac{1}{4L} \sum_{k=0}^{2L-1} \hat{f}_k \left(e^{\frac{2i\pi kn}{2L}} + e^{-\frac{2i\pi kn}{2L}} e^{-\frac{2i\pi k}{2L}} \right) = \\ &= \frac{1}{4L} \sum_{k=0}^{2L-1} \hat{f}_k e^{-\frac{i\pi k}{2L}} \left(e^{\frac{2i\pi k}{2L} \left(n + \frac{1}{2}\right)} + e^{-\frac{2i\pi k}{2L} \left(n + \frac{1}{2}\right)} \right) = \frac{1}{4L} \sum_{k=0}^{2L-1} \hat{f}_k e^{-\frac{i\pi k}{2L}} \cdot 2 \cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) \\ &= \frac{1}{2L} \left(\sum_{k=0}^{L-1} \hat{f}_k e^{-\frac{i\pi k}{2L}} \cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) + \sum_{k=1}^{L-1} \hat{f}_{2L-k} e^{-\frac{i\pi(2L-k)}{2L}} \cos\left(\frac{\pi(2L-k)}{L} \left(n + \frac{1}{2}\right)\right) \right) \end{aligned}$$

Next, we simplify the expression for the second sum:

$$\begin{aligned} e^{-\frac{i\pi(2L-k)}{2L}} \cos\left(\frac{\pi(2L-k)}{L} \left(n + \frac{1}{2}\right)\right) &= e^{-i\pi} e^{\frac{i\pi k}{2L}} \cos\left(2\pi n + \pi - \frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) = \\ &= -e^{\frac{i\pi k}{2L}} \cdot \left(-\cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) \right) = e^{\frac{i\pi k}{2L}} \cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) \end{aligned}$$

Using this result, we obtain:

$$\begin{aligned}
f_n &= \frac{1}{2L} \left(\hat{f}_0 + \sum_{k=0}^{L-1} \hat{f}_k e^{-i\frac{\pi k}{2L}} \cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) + \sum_{k=1}^{L-1} \hat{f}_{2L-k} e^{i\frac{\pi k}{2L}} \cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right) + \hat{f}_{2L} \right) \\
&= \frac{1}{2L} \sum_{k=0}^{L-1} \left(\hat{f}_k e^{-i\frac{\pi k}{2L}} + \hat{f}_{2L-k} e^{i\frac{\pi k}{2L}} \right) \cos\left(\frac{\pi k}{L} \left(n + \frac{1}{2}\right)\right)
\end{aligned}$$

Finally, the coefficients a_k are given by:

$$a_k = \frac{\hat{f}_k e^{-i\frac{\pi k}{2L}} + \hat{f}_{2L-k} e^{i\frac{\pi k}{2L}}}{2L}$$

Question 1.2

To demonstrate that this basis is orthogonal, we compute the inner product $\langle e_k, e_j \rangle$ for two basis vectors. We analyze two cases: when $k = j$ (self-product) and when $k \neq j$ (cross-product).

The inner product is defined as:

$$\begin{aligned}
\langle e_k, e_j \rangle &= \sum_{n=0}^{L-1} \lambda_k \sqrt{\frac{2}{L}} \cos\left(\frac{k\pi}{L} \left(n + \frac{1}{2}\right)\right) \lambda_j \sqrt{\frac{2}{L}} \cos\left(\frac{j\pi}{L} \left(n + \frac{1}{2}\right)\right) = \\
&= \frac{\lambda_k \lambda_j 2}{L} \sum_{n=0}^{L-1} \frac{\cos\left(\frac{\pi}{L}(k+j) \left(n + \frac{1}{2}\right)\right) + \cos\left(\frac{\pi}{L}(k-j) \left(n + \frac{1}{2}\right)\right)}{2} \\
&= \frac{\lambda_k \lambda_j}{L} \left(\sum_{n=0}^{L-1} \cos\left(\frac{\pi}{L}(k+j) \left(n + \frac{1}{2}\right)\right) + \sum_{n=0}^{L-1} \cos\left(\frac{\pi}{L}(k-j) \left(n + \frac{1}{2}\right)\right) \right)
\end{aligned}$$

If $k \neq j$, summing both cosine terms will yield 0 because they form a symmetric function.

If $k = j \neq 0$:

$$\langle e_k, e_k \rangle = \frac{\lambda_k^2}{L} \left(\sum_{n=0}^{L-1} \cos\left(\frac{2k\pi}{L} \left(n + \frac{1}{2}\right)\right) + L \right) = \frac{1}{L} (0 + L) = \frac{1}{L} \cdot L = 1$$

If $k = j = 0$:

$$\langle e_0, e_0 \rangle = \frac{\lambda_0^2}{L} (L + L) = \frac{1}{2L} \cdot 2L = 1$$

Question 1.3

Decomposition of f_n in terms of the basis:

$$f_n = \sum_{k=0}^{L-1} a_k e_k(n)$$

Since $\{e_k\}$ is orthonormal, we can calculate the coefficients a_k as:

$$a_k = \langle f, e_k \rangle = \sum_{n=0}^{L-1} f_n e_k(n)$$

Substituting the definition of $e_k(n)$:

$$a_k = \lambda_k \sqrt{\frac{2}{L}} \sum_{n=0}^{L-1} f_n \cos\left(\frac{k\pi}{L} \left(n + \frac{1}{2}\right)\right)$$

Question 1.4

The program to decompose a signal f in the discrete cosine transform:

```

1 def dct_transform(f):
2     L = len(f)
3     a_k = np.zeros(L) # Array to store DCT coefficients
4     for k in range(L):
5         # Definition of _k
6         if k == 0:
7             lambda_k = 1 / np.sqrt(2)
8         else:
9             lambda_k = 1
10        # Compute the coefficient a_k
11        for n in range(L):
12            a_k[k] += lambda_k * np.sqrt(2 / L) * f[n] * np.cos((np.pi * k / L) * (n + 0.5))
13    return a_k

```

Question 1.5

Suppose we have two vectors:

- $\mathbf{u} \in \mathbb{R}^m$, written as $\mathbf{u} = [u_1, u_2, \dots, u_m]$,
- $\mathbf{v} \in \mathbb{R}^n$, written as $\mathbf{v} = [v_1, v_2, \dots, v_n]$.

Then their tensor product $\mathbf{u} \otimes \mathbf{v}$ is a new vector of size $m \times n$, whose elements are defined as:

$$\mathbf{u} \otimes \mathbf{v} = \begin{bmatrix} u_1 v_1 & u_1 v_2 & \cdots & u_1 v_n \\ u_2 v_1 & u_2 v_2 & \cdots & u_2 v_n \\ \vdots & \vdots & \ddots & \vdots \\ u_m v_1 & u_m v_2 & \cdots & u_m v_n \end{bmatrix}.$$

Thus, using a tensor product of 1D basis we get an orthonormal basis for images of size $L \times L$:

$$e_{k_1, k_2}(n_1, n_2) = \lambda_{k_1} \lambda_{k_2} \frac{2}{L} \cos\left(\frac{k_1 \pi}{L} \left(n_1 + \frac{1}{2}\right)\right) \cos\left(\frac{k_2 \pi}{L} \left(n_2 + \frac{1}{2}\right)\right),$$

where:

- n_1, n_2 are the coordinates of a pixel in the image,
- k_1, k_2 are the indices of the basis functions.

The image function $f(n_1, n_2)$ can be represented as:

$$f(n_1, n_2) = \sum_{k_1=0}^{L-1} \sum_{k_2=0}^{L-1} a_{k_1, k_2} e_{k_1, k_2}(n_1, n_2),$$

where the coefficients a_{k_1, k_2} could be computed as:

$$a_{k_1, k_2} = \sum_{n_1=0}^{L-1} \sum_{n_2=0}^{L-1} f(n_1, n_2) e_{k_1, k_2}(n_1, n_2).$$

Exercise 2. Introduction to JPEG

Question 2.1

We write a program to read an image named `cameraman.png` and compute its DCT transform. The orthonormal basis at each point (n_1, n_2) is computed using an auxiliary function.

In our case, the original matrix contains 256×256 elements. When we split it into blocks of 8×8 , we obtain:

$$\left(\frac{256}{8}\right)^2 = 32 \times 32 = 1024$$

matrices. These blocks are stored as a 3D array with dimensions $(1024, 8, 8)$.

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 def DCT(f):
5     ak = np.zeros_like(f)
6     L = len(f)
7
8     for k1 in range(L):
9         lambda_k1 = 1 / np.sqrt(2) if k1 == 0 else 1
10        for k2 in range(L):
11            lambda_k2 = 1 / np.sqrt(2) if k1 == 0 else 1
12            for n1 in range(L):
13                for n2 in range(L):
14                    ak[k1,k2] += lambda_k1 * lambda_k2 * 2 / L * f[n1,n2] * np.cos((np.pi * k1 / L) *
15                    ↪ (n1 + 0.5)) * np.cos((np.pi * k2 / L) * (n2 + 0.5))
16
17    return ak
18
19 # Processing of image
20 image = plt.imread('cameraman.png') # Load the image (convert to grayscale if necessary)
21 image_array = np.array(image) # Convert the image to a NumPy array
22
23 # Block size
24 block_size = 8
25
26 # Create a list to store all the blocks
27 blocks = []
28
29 # Split the image into 8x8 blocks
30 for i in range(0, 256, block_size):
31     for j in range(0, 256, block_size):
32         block = image_array[i:i + block_size, j:j + block_size]
33         blocks.append(block)
34
35 blocks_array = np.array(blocks)
36
37 # Performing the DCT on each block
38 blocks_array_DCT = np.zeros(np.shape(blocks_array))
39 for i in range(blocks_array.shape[0]):
40     blocks_array_DCT[i] = DCT(blocks_array[i])
```

Question 2.2

The coefficients a_k for each block were computed using the formula derived in **Question 1.5**.

As a result of the program described in **Question 2.1**, we obtain the following coefficients for the first 8×8 matrix out of the 1024 matrices.

$$\begin{bmatrix} 4.9211 & -0.0055 & 0.0018 & 0.0013 & 0.0066 & 0.0029 & 0.0071 & -0.0014 \\ 0.0212 & -0.0028 & 0.0021 & 0.0028 & -0.0077 & 0.0008 & -0.0090 & 0.0157 \\ 0.0113 & 0.0020 & 0.0132 & -0.0057 & -0.0093 & -0.0062 & -0.0019 & -0.0026 \\ -0.0277 & 0.0003 & -0.0124 & -0.0018 & 0.0036 & -0.0036 & 0.0021 & -0.0011 \\ 0.0187 & -0.0017 & 0.0013 & 0.0024 & 0.0034 & 0.0011 & -0.0085 & 0.0021 \\ -0.0115 & 0.0036 & 0.0044 & -0.0056 & -0.0100 & 0.0002 & 0.0015 & 0.0056 \\ 0.0111 & -0.0029 & 0.0040 & -0.0054 & -0.0084 & -0.0213 & -0.0083 & -0.0030 \\ -0.0159 & 0.0039 & -0.0050 & -0.0080 & -0.0050 & -0.0016 & -0.0149 & 0.0123 \end{bmatrix}$$

Question 2.3

We compute the quantization matrix $Q(i,j) = 1 + q(1 + i + j)$ with the use of auxiliary function:

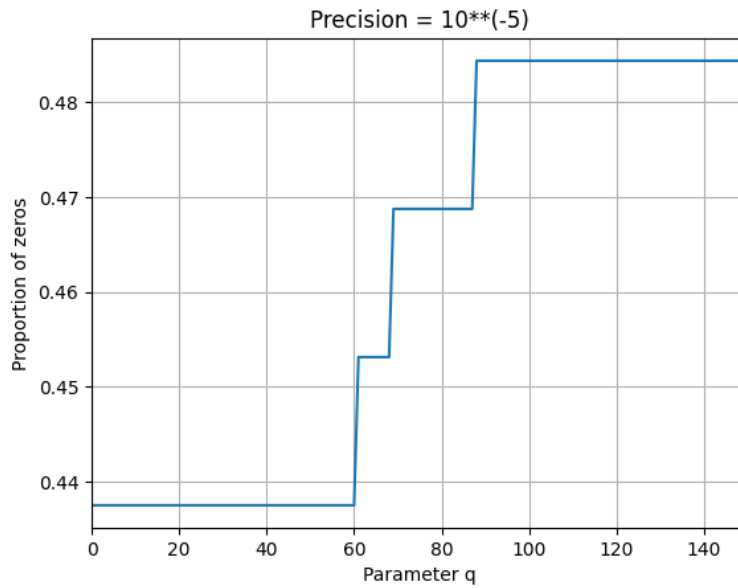
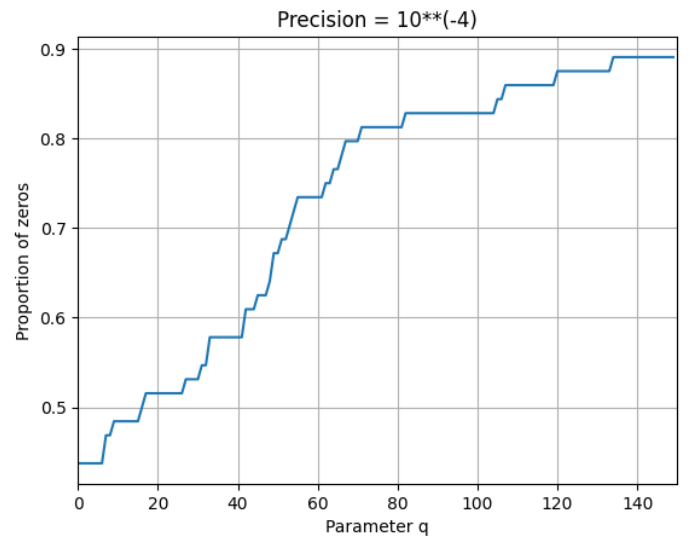

```
1 def quan_matrix(q,block_size):
2     Q = np.zeros((block_size,block_size))
3     for i in range(block_size):
4         for j in range(block_size):
5             Q[i,j] = 1 + q * (1 + i + j)
6     return Q
```

Question 2.4

We plot the proportion of zero coefficients as a function of q . The most important parameter here is precision, because the matrix does not contain exact zeros but rather very small values. If a value is smaller than the precision, we assume it is equal to zero.

There are three graphs, corresponding to precision values of 10^{-3} , 10^{-4} , and 10^{-5} , respectively.

```
1 # Define the maximum range of the quantization parameter q
2 ax_len = 150
3 proportion_of_zeros = np.zeros(ax_len)
4
5 # Define the precision threshold below which values are considered zero
6 precision = 10**(-4)
7
8 blocks_array_DCT_quan = np.zeros(np.shape(blocks_array))
9
10 # Loop through different quantization levels (q values)
11 for q in range(0, ax_len):
12     # Generate the quantization matrix for the current q value
13     Q = quan_matrix(q, block_size)
14
15     for i in range(blocks_array_DCT.shape[0]):
16         # Quantize the DCT coefficients by dividing by the quantization matrix Q
17         blocks_array_DCT_quan[i] = blocks_array_DCT[i] / Q
18
19     # Calculate the proportion of coefficients that are effectively zero (less than
20     ↪ precision)
21     proportion_of_zeros[q] = np.sum(blocks_array_DCT_quan[i] < precision) / (block_size**2)
```



We observe that with the increase of q , the number of zero elements rises significantly. For precision 10^{-3} , the proportion of zeros grows rapidly, and at a compression ratio of approximately 20, the number of zeros exceeds 90%.

For precision 10^{-4} , the graph grows more gradually, reaching about 90% zeros at a compression ratio of 120. However, for precision 10^{-5} , the quantization is insufficient to make the coefficients smaller than this threshold.

Question 2.5

In the DCT matrix, low-frequency coefficients, located in the top-left corner, typically have larger magnitudes, while high-frequency coefficients, in the bottom-right, are much smaller. During quantization, a quantization matrix is applied to reduce the precision of the coefficients, where higher-frequency coefficients are more aggressively quantized due to their smaller magnitudes. This often results in many high-frequency coefficients being rounded to zero. The zigzag order organizes the coefficients such that the most significant, low-frequency coefficients appear first, followed by the high-frequency coefficients, which are more likely to be zero.